

deepseek-ai/**DeepSeek-V4-Flash-0731**

like 576

Follow DeepSeek 139k

Text Generation Transformers Safetensors deepseek_v4 8-bit precision fp8

arxiv:2606.19348 License: mit

Deploy Copy to bucket **NEW** Use this model

Model card Files xet Community 13

DeepSeek-V4-Flash-0731



DeepSeek Homepage Chat DeepSeek V4

Hugging Face DeepSeek AI Twitter deepseek ai

License MIT

Technical Report

Introduction

DeepSeek-V4-Flash-0731 is the official release of **DeepSeek-V4-Flash**, superseding the preview version, with substantially enhanced agentic capabilities. It has the same model structure as DeepSeek-V4-Flash-DSpark, i.e. it comes with a speculative decoding module attached.

DeepSeek-V4-Flash-0731 outperforms DeepSeek-V4-Pro (Preview) on benchmarks listed below

Downloads last month

Safetensors

Model size 304B params

Tensor type BF16 · I64 · F32 · F8_E4M3 · I8

Files info

Inference Providers **NEW**

Text Generation

This model isn't deployed by any Inference Provider.

11 Ask for provider support

Model tree for deepseek-ai/DeepS...

Quantizations 6 models

Space using deepseek-ai/DeepS... 1

victor/DeepSeek-V4-Flash-0731-free...

Collection including deepseek-ai/...

despite its far smaller activated parameter count, and is broadly competitive with the strongest proprietary models available.

Benchmark	DeepSeek-V4-Flash-0731	DeepSeek-V4-Flash (Preview)	DeepSeek-V4-Pro (Preview)
Terminal Bench 2.1	82.7	61.8	72.1
NL2Repo	54.2	39.4	38.5
Cybergym	76.7	38.7	52.7
DeepSWE	54.4	7.3	12.8
Toolathlon-Verified	70.3	49.7	55.9
Agents' Last Exam	25.2	15.8	16.5
AutomationBench Public	25.1	10.8	12.8
DSBench-FullStack †	68.7	37.0	41.8
DSBench-Hard †	59.6	25.8	31.1

Notes:

1. For the Code Agent tasks among the public benchmarks above, DeepSeek-V4-Flash-0731 is evaluated with the minimal mode of DeepSeek Harness (to be released) as the agent framework, using the max reasoning effort level with temperature = 1.0, top_p = 0.95.

DeepSeek-V4

Collection

7 items • Updated about 2 h... • 767

 **Paper for** deepseek-ai/DeepSeek-...

DeepSeek-V4: Towards Highly Effici...

Paper • 2606.19348 • Publi... • 18

2. † DSBench-FullStack is an internal full-stack development test set; DSBench-Hard is an internal test set of difficult coding-agent problems.

Chat Template

This release does not include a Jinja-format chat template. Instead, we provide a dedicated encoding folder with Python scripts and test cases demonstrating how to encode messages in OpenAI-compatible format into input strings for the model, and how to parse the model's text output. Please refer to the [encoding](#) folder for full documentation.

The `reasoning_effort` parameter now supports three levels — `low`, `high`, and `max` — which control how much deliberation the model spends before answering.

A brief example:

```
from encoding_dsv4 import encode_messages, parse_messages

messages = [
    {"role": "user", "content": "hello"},
    {"role": "assistant", "content": "Hello!"},
    {"role": "user", "content": "1+1=?"}
]

# messages -> string
prompt = encode_messages(messages, thinking_effort="low")

# string -> tokens
import transformers
tokenizer = transformers.AutoTokenizer.from_pretrained("gpt-4o")
tokens = tokenizer.encode(prompt)
```

How to Run with vLLM

DSpark speculative decoding is enabled with a single flag — add `--speculative-config` with method: `dspark` to your vLLM launch command:

```
--speculative-config  
'{"method":"dspark","num_speculative_tokens":7  
, "draft_sample_method":"greedy"}'
```

For example, the command below serves the model with vLLM on a single 4×GB300 node. See the [vLLM recipe](#) for detailed instructions and other hardware configurations.

```
vllm serve deepseek-ai/DeepSeek-V4-Flash-073:  
  --trust-remote-code --kv-cache-dtype fp8 --  
  --data-parallel-size 4 --enable-expert-par  
  --moe-backend deep_gemm_mega_moe \  
  --attention-config '{"use_fp4_indexer_cache":  
  --speculative-config '{"method":"dspark", "i
```

How to Run Locally

Please refer to the [inference](#) folder for detailed instructions on running DeepSeek-V4 locally, including model weight conversion and interactive chat demos.

For local deployment, we recommend setting the sampling parameters to `temperature = 1.0`, with `top_p = 0.95` for agentic scenarios and `top_p = 1.0` otherwise. For the high and max reasoning effort levels, we recommend a maximum output length of **384K** tokens.

License

This repository and the model weights are licensed under the [MIT License](#).

Citation

```
@misc{deepseekai2026deepseekv4,  
  title={DeepSeek-V4: Towards Highly Eff:  
  author={DeepSeek-AI},  
  year={2026},  
}
```

Contact

If you have any questions, please raise an issue or contact us at service@deepseek.com.