

 Dogacel / **specdrift-gpt-oss-120b-eagle3** 

 like 0

 Safetensors

 Dogacel/nemotron-post-training-v2-gpt-oss-120b-regen

 English

llama

 arxiv:2605.09992

 License: cc-by-4.0



→ Copy to bucket **NEW**

 **Model card**

 Files

 xet

 Community

Model Card for Model ID

EAGLE-3 drafter model for GPT-oss-120b. This model is released as a part of *Attention Drift: What Speculative Decoding Models Learn* paper. It has several minor architectural differences from the original EAGLE: Drafter hidden state is captured *after* the norm, additional norm injected before FC.

Model Details

Model Sources [optional]

- **Repository:** [Dogacel/SpecDrift](#)
- **Paper:** <https://arxiv.org/abs/2605.09992>

Uses

We recommend using SGLang to run the model,

```
export SGLANG_ENABLE_SPEC_V2=1
```

```
python -m sglang.launch_server \
  --model-path openai/gpt-oss-120b \
  --speculative-algorithm EAGLE3 \
```

Downloads last month
433



 **Safetensors** 

Model size 0.4B params

Tensor type I64 · BF16 · BOOL

 Files info

 **Inference Providers** **NEW**

This model isn't deployed by any Inference Provider.

 Ask for provider support

 **Model tree for Dogacel/specdrift...**

Base model [openai/gpt-oss-120b](#)

 **Finetuned** (110) [this model](#)

 **Dataset used to train Dogacel/spe...**

 Dogacel/nemotron-post-train...

 Viewer • Update... •  930k •  109

```
--speculative-draft-model-path "Dogacel/!
--speculative-num-steps 3 \
--speculative-eagle-topk 1 \
--speculative-num-draft-tokens 4 \
--speculative-draft-sliding-window 2048 `
--port 30000 \
--dp-size 1 --tp-size 1 \
--max-running-requests 64 \
--cuda-graph-max-bs 64 \
--attention-backend fa3 \
--trust-remote-code \
--mem-fraction-static 0.95 --dtype bfloat16
```

Training Details

Training Data

This model is trained on Nemoron Post Training V2 dataset, answers regenerated using gpt-oss-120b.

Dataset publicly available at:

<https://huggingface.co/datasets/Dogacel/nemotron-post-training-v2-gpt-oss-120b-regen>

Training Procedure

We've trained our model using SpecForge on 8xH200 within 10 hours.

- **LR:** 1e-4 (warmup 0.2, cosine)
- **Epochs:** 2
- **Batch Size:** 2 (Effective 2x8=16)
- **Max Length:** 4096
- **TTT:** 4

Evaluation

📦 **Collection including** Dogacel/spec...

SpecDrift 📦 Collection

Models releas... • 2 items • U... • △ 2

📄 **Paper for** Dogacel/specdrift-gpt...

Attention Drift: What Autoregressiv...

📄 Paper • 2605.09992 • Publis... • △ 1

Evaluation has run on: MT-Bench, 80 prompts, max tokens 2048, temperature 0.7

Scripts available at [SpecForge](#).

gpt-oss-120b · EAGLE3 (1-3-1-4) on H100

Throughput (AL) and Δ vs Baseline

BS	Baseline	D-Flash	Ours (1-3-1-4)
1	212	231 (2.47) +8.5%	285 (2.41) +34.1%
8	795	905 (2.47) +13.9%	1044 (2.41) +31.4%
64	1620	1730 (2.42) +6.8%	2339 (2.43) +44.4%

Throughput in tok/s.

Per-Category Throughput

BS=1

Category	Baseline	Ours	AL	Δ
Writing	178.55	214.42	2.340	+20.1%
Roleplay	192.61	296.90	2.270	+54.1%
Reasoning	130.13	168.49	2.393	+29.5%
Math	84.26	115.32	3.101	+36.9%
Coding	329.93	463.48	2.781	+40.5%

Category	Baseline	Ours	AL	Δ
Extraction	98.13	123.12	2.611	+25.5%
STEM	335.91	445.40	2.339	+32.6%
Humanities	349.63	451.13	2.131	+29.0%

BS=8

Category	Baseline	Ours	AL	Δ
Writing	513.57	734.62	2.272	+43.0%
Roleplay	856.44	1149.83	2.250	+34.3%
Reasoning	481.93	679.52	2.370	+41.0%
Math	362.71	443.77	3.160	+22.4%
Coding	1278.03	1646.52	2.765	+28.8%
Extraction	388.07	486.76	2.658	+25.4%
STEM	1240.15	1613.25	2.325	+30.1%
Humanities	1238.40	1600.98	2.172	+29.3%

BS=64

Category	Baseline	Ours	AL	Δ
Writing	1251.72	1855.41	2.401	+48.2%
Roleplay	1477.10	2389.35	2.238	+61.8%
Reasoning	947.08	1488.27	2.380	+57.1%
Math	641.84	997.58	3.095	+55.4%

Category	Baseline	Ours	AL	Δ
Coding	2591.83	3607.85	2.803	+39.2%
Extraction	784.31	1139.90	2.733	+45.3%
STEM	2634.67	3753.09	2.315	+42.5%
Humanities	2630.41	3479.54	2.193	+32.3%

Citation

BibTeX:

```
@misc{eldenk2026attentiondrift,
  title={Attention Drift: What Autoregressive Models Can Do},
  author={Doğaç Eldenk and Payal Mohapatra},
  year={2026},
  eprint={2605.09992},
  archivePrefix={arXiv},
  primaryClass={cs.LG},
  url={https://arxiv.org/abs/2605.09992}
}
```

Acknowledgements

We would like to thank fal and Lambda for their