



hexgrad/Kokoro-82M

like 6.19k



Text-to-Speech



English

doi:10.57967/hf/4329

arxiv:2306.07691

arxiv:2203.02395



License: apache-2.0



Deploy

Copy to bucket

NEW

Use this model



Model card



Files



xet



Kokoro is an open-weight TTS model with 82 million parameters. Despite its lightweight architecture, it delivers comparable quality to larger models while being significantly faster and more cost-efficient. With Apache-licensed weights, Kokoro can be deployed anywhere from production environments to personal projects.

0:00 / 0:20

GitHub: <https://github.com/hexgrad/kokoro>Demo: <https://hf.co/spaces/hexgrad/Kokoro-TTS>

As of April 2025, the market rate of Kokoro served over API is **under \$1 per million characters of text input**, or under \$0.06 per hour of audio output. (On average, 1000 characters of input is about 1 minute of output.) Sources: [ArtificialAnalysis/Replicate at 65 cents per M chars](#) and [DeepInfra at 80 cents per M chars](#).

This is an Apache-licensed model, and Kokoro has been deployed in numerous projects and commercial APIs. We welcome the deployment of the model in real use cases.

Fake websites like [kokorottsai_com](#) (snapshot: <https://archive.ph/nRRnk>) and [kokorotts_net](#) (snapshot: <https://archive.ph/60opa>) are likely scams masquerading under the banner of a popular model.

Any website containing "kokoro" in its root domain (e.g. [kokorottsai_com](#), [kokorotts_net](#)) is **NOT owned by and NOT affiliated with this model page or its author**, and attempts to imply otherwise are red flags.

Downloads last month
10,756,374

Inference Provider

Text-to-Speech

Your prompt here

Run 15,000+ Models Instantly

Inference Providers let you run inference on thousands of models served by our partners using a simple, unified, OpenAI-compatible serverless API ([Learn more](#)).

hexgrad/Kokoro-82M is supported by the following Inference Providers:



View API Code

Dismiss

Generate

</> View Code

Maximize

Model tree for hexgrad/Kokoro-82M

Base model y14579/StyleTTS2-LJSpee...

Finetuned (39) [this model](#)

Adapters 9 models

Finetunes 25 models

Quantizations 39 models

Spaces using hexgrad/Kokoro... 100

hexgrad/Kokoro-TTS

TTS-AGI/TTS-Arena-V2

Pendrokar/TTS-Spaces-Arena

Remsky/Kokoro-TTS-Zero

Nymbo/Tools

ravediamond/baby-reachy-mini-co...

Inferless/Open-Source-TTS-Gallery

TTS-AGI/TTS-Arena + 92 Spaces

Papers for hexgrad/Kokoro-82M

StyleTTS 2: Towards Human-Level T...

- [Releases](#)
- [Usage](#)
- [EVAL.md](#) 
- [SAMPLES.md](#) 
- [VOICES.md](#) 
- [Model Facts](#)
- [Training Details](#)
- [Creative Commons Attribution](#)
- [Acknowledgements](#)

Releases

Model	Published	Training Data	Langs & Voices	SHA256
v1.0	2025 Jan 27	Few hundred hrs	8 & 54	496dba11
v0.19	2024 Dec 25	<100 hrs	1 & 10	3b0c392f

Training Costs	v0.19	v1.0	Total
in A100 80GB GPU hours	500	500	1000
average hourly rate	\$0.80/h	\$1.20/h	\$1/h
in USD	\$400	\$600	\$1000

Usage

You can run this basic cell on [Google Colab](#). [Listen to samples](#). For more languages and details, see [Advanced Usage](#).

```
!pip install -q kokoro>=0.9.2 soundfile
!apt-get -qq -y install espeak-ng > /dev/nul
from kokoro import KPipeline
from IPython.display import display, Audio
import soundfile as sf
import torch
pipeline = KPipeline(lang_code='a')
text = '''
[Kokoro](/k'Okəɹ0/) is an open-weight TTS mo
'''
```

```
generator = pipeline(text, voice='af_hearth')
for i, (gs, ps, audio) in enumerate(generator):
    print(i, gs, ps)
    display(Audio(data=audio, rate=24000, au
sf.write(f'{i}.wav', audio, 24000)
```

Under the hood, kokoro uses [misaki](#), a G2P library at <https://github.com/hexgrad/misaki>

Model Facts

Architecture:

- StyleTTS 2: <https://arxiv.org/abs/2306.07691>
- ISTFTNet: <https://arxiv.org/abs/2203.02395>
- Decoder only: no diffusion, no encoder release

Architected by: Li et al @ <https://github.com/yl4579/StyleTTS2>

Trained by: @izzzn on Discord

Languages: Multiple

Model SHA256 Hash:

496dba118d1a58f5f3db2efc88dbdc216e0483fc89fe6e
47ee1f2c53f18ad1e4

Training Details

Data: Kokoro was trained exclusively on **permissive/non-copyrighted audio data** and IPA phoneme labels. Examples of permissive/non-copyrighted audio include:

- Public domain audio
 - Audio licensed under Apache, MIT, etc
 - Synthetic audio^[1] generated by closed^[2] TTS models from large providers
- [1] https://copyright.gov/ai/ai_policy_guidance.pdf
- [2] No synthetic audio from open TTS models or "custom voice clones"

Total Dataset Size: A few hundred hours of audio

Total Training Cost: About \$1000 for 1000 hours of A100 80GB vRAM

Creative Commons Attribution

The following CC BY audio was part of the dataset used to train Kokoro v1.0.

Audio Data	Duration Used	License	Added to Training Set After
Koniwa tnc	<1h	CC BY 3.0	v0.19 / 22 Nov 2024
SIWIS	<11h	CC BY 4.0	v0.19 / 22 Nov 2024

Acknowledgements

- 🛠️ [@yl4579](#) for architecting StyleTTS 2.
- 🏆 [@Pendrokar](#) for adding Kokoro as a contender in the TTS Spaces Arena.
- 📊 Thank you to everyone who contributed synthetic training data.
- ❤️ Special thanks to all compute sponsors.
- 🌐 Discord server: <https://discord.gg/QuGxSWBfQy>.
- 🍃 Kokoro is a Japanese word that translates to "heart" or "spirit". It is also the name of an [AI in the Terminator franchise](#).

