

incoai/**GLM-5.3-Flash-DFlash2**
 like 97

Follow  Inco AI 222

 Text Generation

 Transformers

 Safetensors

qwen3

>> dflash

dflash2

speculative-decoding

block-diffusion

draft-model

sglang

 text-generation-inference

 License: cc-by-nc-nd-4.0

 Deploy ▾

→ Copy to bucket **NEW**

Use this model ▾

 **Model card**

Files

 xet

 Community 4

GLM-5.3-Flash-DFlash2

[Blog](#) | [GitHub](#)

This repository contains the DFlash 2 draft model for [zai-org/GLM-5.3-Flash](#). It is not a standalone language model: it runs inside a speculative decoding server and drafts tokens for the target model to verify.

DFlash 2 is a block-diffusion drafter for speculative decoding. It predicts a whole block of tokens in a single pass and keeps the top candidates at every position. A lightweight selector then traces one coherent path through them. Two-tap dynamic convolutions in the backbone keep the draft from decaying toward the end of the block. Decoding is lossless: greedy output matches the target model exactly, and sampling preserves its distribution.

Downloads last month

9,190

 **Safetensors** ⓘ

Model size 1B params

Tensor type BF16 [Files info](#)
 **Inference Providers** **NEW**
 Text Generation

This model isn't deployed by any Inference Provider.

 Ask for provider support


 **Model tree for** incoai/GLM-5.3-F1...

Base model [zai-org/GLM-5.3-Flash](#)
 **Finetuned** (12) [this model](#)

Quantizations ..    4 models

 **Collection including** incoai/GLM-5...

DFlash 2  Collection



Quick Start

Serve with [SGLang](#):

```
pip install "sglang[all] @ git+https://github.com
```

```
sglang serve \
  --model-path zai-org/GLM-5.3-Flash \
  --trust-remote-code \
  --speculative-algorithm DFLASH \
  --speculative-draft-model-path incoai/GLM-5.3-Flash \
  --speculative-draft-attention-backend fa4
```

See the [blog post](#) for more details.

Evaluation

Benchmark numbers for this checkpoint are being re-measured and will be published here once final.

License

This model is released under [CC BY-NC-ND 4.0](#) for research and evaluation. For commercial licensing, contact contact@inco.ai.

Citation

If you find DFlash 2 useful, please cite:

```
@misc{inco2026dflash2,  
  title = {{DFlash 2: Keep Drafting Parallel}},  
  author = {{Inco AI}},  
  year = {2026},  
  month = {August},  
  url = {https://inco.ai/blog/dflash2/}  
}
```

Please also cite the original DFlash paper:

```
@inproceedings{chen2026dflash,  
  title = {{DFlash: Block Diffusion for LLMs}},  
  author = {Chen, Jian and Liang, Yesheng},  
  booktitle = {International Conference on Machine Learning},  
  year = {2026}  
}
```