

incoai/**GLM-5.3-NVFP4**


like

17

Follow

inco Inco AI

233



Text Generation



Transformers



Safetensors

glm_moe_dsa

GLM-5

ModelOpt

quantized

4-bit precision

FP4

nvfp4

conversational



8-bit precision

modelopt



License: glm-5.3



Deploy

→ Copy to bucket

NEW

Use this model


Model card


Files



xet



Community

GLM-5.3-NVFP4

This repository contains an NVFP4-quantized version of Z.ai's [zai-org/GLM-5.3](#). The weights and activations of the routed-expert linear layers are in NVFP4, and the KV cache is in FP8. The checkpoint uses the NVIDIA Model Optimizer format (quant_method: modelopt) and is served directly by SGLang and vLLM on NVIDIA Blackwell GPUs (SM100+).

- **Base model:** [zai-org/GLM-5.3](#)
- **Quantization:** NVFP4 weights + static NVFP4 activations; FP8 KV cache
- **Size:** 433 GiB
- **License:** [GLM-5.3 License](#)

Quantization Method

This checkpoint was produced with our in-house post-training quantization toolchain and exported in the [NVIDIA Model Optimizer](#) format. Only the weights and activations of the linear operators

Downloads last month

10,412


Safetensors ?

Model size 391B params

Tensor type U8 · BF16 · F32

 Chat template

 Files info

Inference Providers NEW

 Text Generation

This model isn't deployed by any Inference Provider.



Ask for provider support

Model tree for incoai/GLM-5.3-NV...

 Base model [zai-org/GLM-5.3](#)
 Quantized (34) [this model](#)
 Quantizations ..    2 models

within the routed MoE experts are quantized to NVFP4, with static per-tensor activation scales. The KV cache is quantized to FP8 with static unit scales (kv_cache_quant_algo: FP8).

Quick Start

The commands below enable speculative decoding with the [DFlash2](#) draft model (lossless, 7 draft tokens per verification step); remove the `--speculative-*` flags to serve without it.

Serve with [SGLang](#) (main):

```
python3 -m sglang.launch_server \
  --model-path incoai/GLM-5.3-NVFP4 \
  --tp 8 \
  --quantization modelopt_fp4 \
  --reasoning-parser glm45 \
  --tool-call-parser glm47 \
  --chunked-prefill-size 8192 \
  --speculative-algorithm DFLASH \
  --speculative-draft-model-path incoai/GLI
  --speculative-draft-attention-backend tr
```

Or with [vLLM](#) (v0.28.0 or later):

```
vllm serve incoai/GLM-5.3-NVFP4 \
  --tensor-parallel-size 8 \
  --reasoning-parser glm45 \
  --tool-call-parser glm47 \
  --enable-auto-tool-choice \
  --speculative-config '{"method":"dflash"}
```

Accuracy

We compare Z.ai's original FP8 release with this NVFP4 checkpoint. Higher is better.

Precision	GPQA Diamond	AIME 2025	MATH-500	HLE	AA-LCR
FP8	91.1	94.3	95.6	35.9	73.6
NVFP4	91.2	95.1	95.2	35.2	73.0

License

This model is a quantized version of GLM-5.3 and is distributed under Z.ai's [GLM-5.3 License](#), which it inherits from the base model.