

lightseekorg / **kimi-k2.6-eagle3-mla**   like 2 Follow  LightSeek Foundation 24

 Text Generation
  Safetensors
 kimi_k2
 speculative-decoding
 eagle3
 kimi-k2.6
 torchspec
  License: other



 **Model card**
 Files
  xet
  Community

kimi-k2.6-eagle3-mla

Model Overview

kimi-k2.6-eagle3-mla is an Eagle3 MTP draft model with MLA (Multi-Latent Attention) for accelerating inference of [Kimi-K2.6](#), trained with [TorchSpec](#) — an online speculative decoding training framework that runs FSDP training and inference concurrently. If you find this draft model useful, please give our project TorchSpec a star on [GitHub](#).

Why an MLA (Multi-Latent Attention) Draft Model

Compared with an MHA draft model, the MLA variant is a better fit for Kimi-K2.6 deployment:

- Uses less KV cache, which reduces serving memory pressure.
- Matches Kimi-K2.6's MLA architecture, so it fits more naturally into the inference engine's KV-cache handling under different serving scenarios such as PD-Disaggregation.

Downloads last month
1,947



Safetensors

Model size 3B params

Tensor type BF16  Files info

Inference Providers NEW

 Text Generation

This model isn't deployed by any Inference Provider.

 Ask for provider support

Model tree for lightseekorg/kimi...

Base model moonshotai/Kimi-K2.6

 Finetuned (12) [this model](#)

Training Setup

- **Cluster:** 3 nodes × 8× B200 (24 GPUs total)
- **Training:** 1 node (8 GPUs), FSDP
- **Inference:** 2 nodes (16 GPUs), vLLM (TP=8 per node)
- **Continual training:** Initialized from [kimi-k2.5-eagle3-mla](#) checkpoint
- **Iterations:** 9,279 steps
- **Learning rate:** 2e-5, cosine schedule

Performance

The primary metric is **accept_length** — the average number of tokens accepted per speculation step with num_speculative_tokens=3. Higher is better.

Benchmarks were run on vLLM 0.20.0 with 8× B200 GPUs.

Category	Benchmark	N	Accept Length
Dialogue	MTBench	80	2.624
Chinese	CEval	212	2.494
Math	GSM8K	500	2.987
Code	HumanEval	164	3.241
Math	MATH500	500	3.245
Math	AIME	30	2.982
Code	LiveCodeBench	200	2.706

Category	Benchmark	N	Accept Length
Code	SPEED-Bench (coding)	80	3.006

Quick Start

Requirements

- NVIDIA GPU with CUDA 12.0+
- [vLLM](#) >= 0.20.0

Launch Server (vLLM)

```
vllm serve moonshotai/Kimi-K2.6 \
  --tensor-parallel-size 8 \
  --speculative-config '{"model": "lightseer"}' \
  --trust-remote-code
```

Launch Server (SGLang)

MLA Eagle3 draft model is not yet supported in SGLang. Will update once support is available.

Citation

```
@misc{torchspec2026,
  title={TorchSpec: An Online Speculative Decoder},
  url={https://github.com/torchspec-project/torchspec},
  year={2026}}
```