

 poolside/**Laguna-S-2.1-DFlash-FP8** 


 like 7

Follow  Poolside 955

 Safetensors  speculators  laguna  speculative-decoding  dflash

 → Copy to bucket **NEW**

 **Model card**  Files  xet  Community 4



[Use on OpenRouter](#) · [Use on Vercel AI Gateway](#) ·
[Release blog post](#)

poolside/Laguna-S-2.1-DFlash-FP8

DFlash speculator for the FP8 target [poolside/Laguna-S-2.1-FP8](#). The speculator is a 6-layer Laguna-style draft model (BF16); pair it with the FP8 base for lower-latency serving via speculative decoding.

Trained: e0630_rhiemann_baseline SFT, DFlash Stage-2, 15k steps. Recommended serving setting: num_speculative_tokens=7. DFlash upstream support is in progress (vLLM #46853, SGLang #29446, TRT-LLM #15666). Use poolside/Laguna-S-2.1-FP8 as the target model.

Benchmarks

Downloads last month
2,452



 **Safetensors** 

Model size 1B params

Tensor type BF16  Files info

 **Inference Providers** **NEW**

This model isn't deployed by any Inference Provider.

 Ask for provider support

 **Model tree for poolside/Laguna-S...**

Base model [poolside/Laguna-S-2.1](#)

 **Quantized** poolside/Laguna-S-...

 **Finetuned (1)** [this model](#)

Quantizations [1 model](#)

 **Collection including poolside/Lag...**

Laguna S 2.1  Collection

Our most ca... • 12 items • U... •  24

Measured with TP=2, temperature=0, and
num_speculative_tokens=15.

Throughput speedup

Concurrency	GSM8K	MATH-500	HumanEval	MBPP
1	3.179x	2.938x	3.269x	2.380x
4	2.614x	2.423x	2.691x	1.963x
8	2.666x	2.410x	2.803x	1.962x
16	2.618x	2.364x	2.866x	2.031x

Acceptance length

Concurrency	GSM8K	MATH-500	HumanEval	MBPP
1	5.748	5.197	5.889	4.247
4	5.765	5.212	5.882	4.218
8	5.863	5.199	6.094	4.178
16	5.787	5.161	6.144	4.291