

 poolside/**Laguna-S-2.1-DFlash-INT4**


 like 4

Follow  Poolside 955

 Safetensors  speculators  laguna  speculative-decoding  dflash

 → Copy to bucket **NEW**

 **Model card**  Files  xet  Community 5



[Use on OpenRouter](#) · [Use on Vercel AI Gateway](#) ·
[Release blog post](#)


poolside/Laguna-S-2.1-DFlash-INT4

DFlash speculator for the INT4 target
[poolside/Laguna-S-2.1-INT4](#). The speculator is a 6-layer Laguna-style draft model (BF16); pair it with the INT4 base for lower-latency serving via speculative decoding.

Trained: e0630_rhiemann_baseline SFT, DFlash Stage-2, 15k steps. Recommended serving setting: num_speculative_tokens=7. DFlash upstream support is in progress (vLLM #46853, SGLang #29446, TRT-LLM #15666). Use poolside/Laguna-S-2.1-INT4 as the target model.

Benchmarks

Downloads last month
2,273



Safetensors

Model size 1B params

Tensor type BF16  Files info

Inference Providers **NEW**

This model isn't deployed by any Inference Provider.

 Ask for provider support

Model tree for poolside/Laguna-S...

Base model [poolside/Laguna-S-2.1](#)

 Quantized poolside/Laguna-S-...

 Finetuned (1) [this model](#)

Collection including poolside/Lag...

Laguna S 2.1  Collection

Our most ca... • 12 items • U... •  24

Measured with TP=2, temperature=0, and
num_speculative_tokens=15.

Throughput speedup

Concurrency	GSM8K	MATH-500	HumanEval	MBPP
1	3.697x	3.216x	3.776x	2.704x
4	2.815x	2.521x	2.995x	2.145x
8	2.541x	2.247x	2.865x	1.970x
16	2.426x	2.161x	2.629x	1.895x

Acceptance length

Concurrency	GSM8K	MATH-500	HumanEval	MBPP
1	6.273	5.363	6.416	4.532
4	6.111	5.336	6.416	4.498
8	6.133	5.356	6.777	4.565
16	6.145	5.393	6.314	4.507