

 poolside/**Laguna-S-2.1-DFlash-NVFP4**


 like 22

Follow  Poolside 955

 Safetensors speculators laguna speculative-decoding dflash

 → Copy to bucket **NEW**

 **Model card**  Files  xet  Community 5



[Use on OpenRouter](#) · [Use on Vercel AI Gateway](#) ·
[Release blog post](#)

poolside/Laguna-S-2.1-DFlash-NVFP4

DFlash speculator for the NVFP4 target
[poolside/Laguna-S-2.1-NVFP4](#). The speculator is a 6-layer Laguna-style draft model (BF16); pair it with the NVFP4 base for lower-latency serving via speculative decoding.

Trained: e0630_rhiemann_baseline SFT, DFlash Stage-2, 15k steps. Recommended serving setting: num_speculative_tokens=7. DFlash upstream support is in progress (vLLM #46853, SGLang #29446, TRT-LLM #15666). Use poolside/Laguna-S-2.1-NVFP4 as the target model.

Benchmarks

Downloads last month
9,215



Safetensors

Model size 1B params

Tensor type BF16  Files info

Inference Providers **NEW**

This model isn't deployed by any Inference Provider.

 Ask for provider support

Model tree for poolside/Laguna-S...

Base model [poolside/Laguna-S-2.1](#)

 Quantized poolside/Laguna-S-...

 Finetuned (1) [this model](#)

Collection including poolside/Lag...

Laguna S 2.1  Collection

Our most ca... • 12 items • U... •  24

Measured with TP=2, temperature=0, and
num_speculative_tokens=15.

Throughput speedup

Concurrency	GSM8K	MATH-500	HumanEval	MBPP
1	3.324x	2.893x	3.692x	2.426x
4	2.498x	2.174x	2.742x	1.848x
8	2.279x	1.948x	2.634x	1.719x
16	2.302x	1.965x	2.626x	1.731x

Acceptance length

Concurrency	GSM8K	MATH-500	HumanEval	MBPP
1	5.775	4.942	6.438	4.171
4	5.804	4.935	6.298	4.161
8	5.719	4.885	6.537	4.193
16	5.758	4.896	6.412	4.123