

 poolside/**Laguna-S-2.1-DFlash** 

like 12

Follow

 Poolside

955

 Safetensors

laguna

dflash

speculative-decoding

 License: other

⋮

→ Copy to bucket **NEW**


 **Model card**

Files

 xet

 Community

5



[Use on OpenRouter](#) · [Use on Vercel AI Gateway](#) ·

[Release blog post](#)

Laguna-S-2.1-DFlash

DFlash speculator (draft model) for [Laguna S 2.1](#), bf16.

- Architecture: DFlashLagunaForCausalLM (6 sliding-attention layers, block_size 16).
- Shares token embedding + lm_head with the target; draft_vocab_size == vocab_size (no d2t/t2d).
- Loads under vLLM (native laguna_dflash) and TRT-LLM (pytorch DFlash backend) as the draft model in a speculative config.

Usage

Downloads last month
2,381



 **Safetensors** 

Model size 1B params

Tensor type BF16

 Files info

 **Inference Providers** **NEW**

This model isn't deployed by any Inference Provider.



Ask for provider support

 **Model tree for poolside/Laguna-S...**

Base model poolside/Laguna-S-2.1

 **Finetuned (7)** [this model](#)

 **Collection including poolside/Lag...**

Laguna S 2.1  Collection

Our most ca... • 12 items • U... •  24

vLLM

```
vllm serve --model poolside/Laguna-S-2.1 \
  --speculative-config '{"model":"poolside,"
  ...
```

SGLang

```
python -m sglang.launch_server --model-path |
  --speculative-algorithm DFLASH \
  --speculative-draft-model-path poolside/La
  ...
```

llama.cpp (GGUF)

A llama.cpp conversion of this draft model (laguna-s-2.1-DFlash-BF16.gguf) is published in [poolside/Laguna-S-2.1-GGUF](#) alongside the target GGUFs. It embeds the target tokenizer and the DFlash metadata (dflash.decoder_arch = laguna, capture layers, block size), so it works directly as the -md draft model:

```
git clone --branch laguna https://github.com,
cd llama.cpp && cmake -B build && cmake --bu:

./build/bin/llama-server \
  -m laguna-s-2.1-Q4_K_M.gguf \
  -md laguna-s-2.1-DFlash-BF16.gguf \
  --spec-type draft-dflash --spec-draft-n-ma:
  -fa on --jinja --port 8000
```

--spec-draft-n-max is clamped to the trained block size (15 draft tokens + 1).

Requires Poolside's llama.cpp fork, branch [laguna](#).
Upstream llama.cpp ships the generic DFlash framework but not the Laguna decoder contract this draft model needs, and upstream PR [ggml-org/llama.cpp#25165](#) covers the target architecture only.

 System theme

[TOS](#)

[Privacy](#)

[About](#)

[Careers](#)



[Models](#)

[Datasets](#)

[Spaces](#)

[Pricing](#)

[Docs](#)