

 poolside/**Laguna-XS-2.1-DFlash**   like 9 Follow  Poolside 955

 Safetensors speculators laguna speculative-decoding dflash  arxiv:2602.06036  License: openmdw-1.1

 → Copy to bucket **NEW**

 **Model card**  Files  xet  Community 2

poolside/Laguna-XS-2.1-DFlash

DFlash speculator for poolside/Laguna-XS-2.1.

Model Specifications

Base Model	poolside/Laguna-XS-2.1
Draft Architecture	5-layer Llama-style draft model (DFlashLagunaForCausalLM)
Hidden Size	2048
Attention Heads	64 query / 8 key-value, head_dim 128
Sliding Window	512
Gating	Per-head
Proposals	Up to 15 tokens per step
Chat Template	poolside/Laguna-XS-2.1 (use /chat/completions endpoint)
Format	Safetensors (BF16)



 **Safetensors** ⓘ

Model size 0.5B params

Tensor type BF16  Files info

 **Inference Providers** **NEW**

This model isn't deployed by any Inference Provider.

 Ask for provider support

 **Model tree for poolside/Laguna-X...**

Base model poolside/Laguna-XS-2.1

 **Finetuned (3)** [this model](#)

 **Quantizations** [1 model](#)

 **Collection including poolside/Lag...**

Laguna XS 2.1  Collection

Designed for ... • 9 items • U... • △ 19

License	OpenMDW-1.1

DFlash: Block Diffusion for Flash Sp...

 Paper • 2602.06036 • Publi... •  89

Deployment

DFlash speculative-decoding support for Laguna XS 2.1 has **not yet landed upstream**. Integrations are in progress:

- **vLLM** — [vllm-project/vllm#46853](https://github.com/vllm-project/vllm/pull/46853)
- **TRT-LLM** — [NVIDIA/TensorRT-LLM#15666](https://github.com/NVIDIA/TensorRT-LLM/pull/15666)

See the "Speculative decoding (DFlash)" notes on the [Laguna XS 2.1 model card](#) for current status.

Once vLLM support lands, pair this speculator with the base model by adding a `--speculative-config` block to the standard Laguna XS 2.1 serve command:

```
# Support in progress – see vllm-project/vllm
vllm serve poolside/Laguna-XS-2.1 \
  --tool-call-parser poolside_v1 \
  --reasoning-parser poolside_v1 \
  --enable-auto-tool-choice \
  --default-chat-template-kwarg '{"enable_
  --speculative-config '{"model": "poolside
```

SGLang support:

```
sglang serve \
  --model-path poolside/Laguna-XS-2.1 \
  --speculative-algorithm DFLASH \
  --speculative-draft-model-path poolside/
```

Evaluation

Speculative-decoding throughput and mean acceptance length for the (BF16) Laguna XS 2.1 target paired with this speculator (num_speculative_tokens = 15), versus the same model without speculative decoding:

Dataset	Baseline (tok/s/seq)	DFlash (tok/s/seq)	Speedup	Acceptance length
GSM8K v2	80.22	133.78	×1.67	3.55
HumanEval	95.57	252.59	×2.64	4.57
EvalPlus	97.68	192.69	×1.97	4.10
Math	73.63	139.84	×1.90	4.40

License

This model is licensed under the [OpenMDW-1.1 License](#).

Intended and Responsible Use

Laguna-XS-2.1-DFlash is designed for software engineering and agentic coding use cases, and you are responsible for confirming that it is appropriate for your intended application. Laguna-XS-2.1-DFlash is subject to the [OpenMDW-1.1 License](#), and should be used consistently with Poolside's [Acceptable Use Policy](#).

Please report security vulnerabilities or safety concerns to security@poolside.ai.

References

- [DFlash: Block Diffusion for Flash Speculative Decoding](#)

 System theme

[TOS](#)

[Privacy](#)

[About](#)

[Careers](#)



[Models](#)

[Datasets](#)

[Spaces](#)

[Pricing](#)

[Docs](#)