

 poolside/ **Laguna-XS.2-speculator.dflash**

Follow  Poolside 955

 Safetensors

speculators

speculative-decoding

dflash

custom_code

 arxiv:2602.06036

 License: apache-2.0

⋮

→ Copy to bucket NEW

 **Model card**

 Files

 xet

 Community

poolside/Laguna-XS.2-speculator.dflash

This is a DFlash speculator model for [poolside/Laguna-XS.2](#).

Training Details

This model was trained using the [Speculators](#) library on a combination of [Magpie-Align/Magpie-Llama-3.1-Pro-300K-Filtered](#) and the train_sft split of [HuggingFaceH4/ultrachat_200k](#). Responses were regenerated by Laguna-XS.2 (with reasoning).

Model Specifications

Base Model	poolside/Laguna-XS.2
Chat Template	poolside/Laguna-XS.2 (use /chat/completions endpoint)
Format	Safetensors
License	Apache 2.0



 Safetensors ⓘ

Model size0.6B params

Tensor typeI64 · BF16 · BOOL

 Files info

 Inference Providers NEW

This model isn't deployed by any Inference Provider.

 Ask for provider support

 Model tree for poolside/Laguna-X...

Base model poolside/Laguna-XS.2

 Finetuned (24) [this model](#)

 Collection including poolside/Lag...

Laguna XS.2  Collection

Designed for ... • 5 items • U... •  28

Validation Hardware	Nvidia A100

DFlash: Block Diffusion for Flash Sp...

 Paper • 2602.06036 • Publi... •  89

Deployment

```
# Requires vLLM >= 0.21.0
pip install 'vllm>=0.21.0'

# Deploy with speculative decoding
VLLM_USE_DEEP_GEMM=0 vllm serve poolside/Laguna-XS.2-speculative-decoding \
  --tensor-parallel-size 1 \
  --max-model-len 16384 \
  --tool-call-parser poolside_v1 \
  --reasoning-parser poolside_v1 \
  --enable-auto-tool-choice \
  --default-chat-template-kwarg '{"enable_tool_calls": true}' \
  --speculative-config '{
    "model": "poolside/Laguna-XS.2-speculative-decoding",
    "num_speculative_tokens": 7,
    "method": "dflash"
  }'
```

Preliminary Evaluations

Per-position token acceptance rates across
datasets: (with reasoning enabled)

Dataset	Pos 1	Pos 2	Pos 3	Pos 4	Pos 5
HumanEval	74.0%	48.6%	29.9%	17.7%	9.9%
math_reasoning	76.9%	53.2%	34.6%	21.2%	12.1%
qa	68.5%	41.8%	24.8%	14.7%	8.4%

Dataset	Pos 1	Pos 2	Pos 3	Pos 4	Pos 5
question	70.6%	44.1%	26.2%	15.0%	8.4%
rag	71.7%	45.7%	27.6%	16.0%	8.9%
summarization	68.8%	40.8%	22.7%	12.3%	6.5%
translation	70.8%	44.3%	25.0%	13.0%	6.5%
writing	70.9%	44.6%	26.8%	15.8%	9.4%

References

Paper: [DFlash: Block Diffusion for Flash Speculative](#)