



Qwen/Qwen2.5-72B-Instruct

like

936

Follow

Qwen

83.3k



Text Generation



Transformers



Safetensors



English

qwen2

chat

conversational



text-generation-inference



arxiv:2309.00071



arxiv:2407.10671



License: qwen



Deploy

Use this model



Model card



Files



xet



Community

36



Qwen2.5-72B-Instruct

Qwen Chat

Introduction

Qwen2.5 is the latest series of Qwen large language models. For Qwen2.5, we release a number of base language models and instruction-tuned language models ranging from 0.5 to 72 billion parameters. Qwen2.5 brings the following improvements upon Qwen2:

- Significantly **more knowledge** and has greatly improved capabilities in **coding** and **mathematics**, thanks to our specialized expert models in these domains.
- Significant improvements in **instruction following**, **generating long texts** (over 8K tokens), **understanding structured data** (e.g, tables), and **generating structured outputs** especially JSON. **More resilient to the diversity of system prompts**, enhancing role-play implementation and condition-setting for chatbots.
- Long-context Support** up to 128K tokens and can generate up to 8K tokens.
- Multilingual support** for over 29 languages, including Chinese, English, French, Spanish, Portuguese, German, Italian, Russian, Japanese, Korean, Vietnamese, Thai, Arabic, and more.

This repo contains the instruction-tuned 72B Qwen2.5 model, which has the following features:

Downloads last month

437,132



Safetensors

Model size 73B params

Tensor type BF16

Chat template

Files info

Inference Providers

Text Generation

Input a message with Qwen/Qwen2.5-72B-Instruct

Run 15,000+ Models Instantly

Inference Providers let you run inference on thousands of models served by our partners using a simple, unified, OpenAI-compatible serverless API ([Learn more](#)).

Qwen/Qwen2.5-72B-Instruct is supported by the following Inference Providers:



Novita



Featherless AI

View API Code

Dismiss

Your prompt here...

Send

</> View Code

Compare providers

Model tree for Qwen/Qwen2.5-72B-Instruct

Base model Qwen/Qwen2.5-72B-Instruct



Finetuned (66)

[this model](#)

Adapters

64 models



Finetunes

63 models



Merges

15 models



Quantizations



82 models

Spaces using Qwen/Qwen2.5-72B-Instruct 100



lablab-ai-amd-developer-hackathon/...



stvident/sqlab



nensi1311/ScholarEnv

- Type: Causal Language Models
- Training Stage: Pretraining & Post-training
- Architecture: transformers with RoPE, SwiGLU, RMSNorm, and Attention QKV bias
- Number of Parameters: 72.7B
- Number of Paramaters (Non-Embedding): 70.0B
- Number of Layers: 80
- Number of Attention Heads (GQA): 64 for Q and 8 for KV
- Context Length: Full 131,072 tokens and generation 8192 tokens
 - Please refer to [this section](#) for detailed instructions on how to deploy Qwen2.5 for handling long texts.

For more details, please refer to our [blog](#), [GitHub](#), and [Documentation](#).

Requirements

The code of Qwen2.5 has been in the latest Hugging face `transformers` and we advise you to use the latest version of `transformers`.

With `transformers<4.37.0`, you will encounter the following error:

```
KeyError: 'qwen2'
```

Quickstart

Here provides a code snippet with `apply_chat_template` to show you how to load the tokenizer and model and how to generate contents.

```
from transformers import AutoModelForCausalLM

model_name = "Qwen/Qwen2.5-72B-Instruct"

model = AutoModelForCausalLM.from_pretrained(
    model_name,
    torch_dtype="auto",
    device_map="auto"
)

tokenizer = AutoTokenizer.from_pretrained(model_name)
```

 [filbench/filbench-leaderboard](#)

 [Adityahars/Sanskrit-env](#)

 [ReviewGrounder/GradioDemo](#)

 [jsdosanj/SikhLibrarian](#)

 [Shasidharyadavr/netweaver_sre](#)

+ 92 Spaces

Collection including Qwen/Qwen2.5...

Qwen2.5  Collection

Qwen2.5 lang... • 43 items • U. •  718

Papers for Qwen/Qwen2.5-72B-Ins...

Qwen2 Technical Report

 Paper • 2407.10671 • Pub... •  172

YaRN: Efficient Context Window Ex...

 Paper • 2309.00071 • Publi... •  83

```

prompt = "Give me a short introduction to la:
messages = [
    {"role": "system", "content": "You are Q
    {"role": "user", "content": prompt}
]
text = tokenizer.apply_chat_template(
    messages,
    tokenize=False,
    add_generation_prompt=True
)
model_inputs = tokenizer([text], return_tensors=

generated_ids = model.generate(
    **model_inputs,
    max_new_tokens=512
)
generated_ids = [
    output_ids[len(input_ids):] for input_ids

]

response = tokenizer.batch_decode(generated_ids)

```

Processing Long Texts

The current `config.json` is set for context length up to 32,768 tokens. To handle extensive inputs exceeding 32,768 tokens, we utilize [YaRN](#), a technique for enhancing model length extrapolation, ensuring optimal performance on lengthy texts.

For supported frameworks, you could add the following to `config.json` to enable YaRN:

```

{
    ...,
    "rope_scaling": {
        "factor": 4.0,
        "original_max_position_embeddings": 32768,
        "type": "yarn"
    }
}

```

For deployment, we recommend using vLLM. Please refer to our [Documentation](#) for usage if you are not familiar with vLLM. Presently, vLLM only supports static YARN, which means the scaling factor remains constant regardless of input length, **potentially impacting performance on shorter texts**. We advise adding the `rope_scaling` configuration only when processing long contexts is required.

