

RedHatAI/GLM-5.3-Flash-NVFP4

 like

10

Follow

 Red Hat AI

3.03k



Image-Text-to-Text



Transformers



Safetensors

glm5_next

vllm

llm-compressor

nvfp4

fp4

conversational



8-bit precision

compressed-tensors


 Deploy

→ Copy to bucket **NEW**

Use this model


Model card


Files



Community 2

RedHatAI/GLM-5.3-Flash-NVFP4

This model is a quantized version of [zai-org/GLM-5.3-Flash](#).

Model Optimizations

This model was obtained by quantizing the weights of [zai-org/GLM-5.3-Flash](#) to NVFP4, ready for inference with vLLM.

Weights are quantized to FP4 with a group size of 16, and activations are quantized to FP4 with local per-group scaling. Only the weights and activations of the linear operators within transformer blocks are quantized using [LLM Compressor](#). Vision tower, embedding, and output head layers are kept in their original precision.

vLLM Serving

```
docker run --gpus all \
  --privileged --ipc=host -p 8000:8000 \
  -v ~/.cache/huggingface:/root/.cache/huggingface \
  -e VLLM_ENGINE_READY_TIMEOUT_S=3600 \
  vllm/vllm-openai:glm53-flash RedHatAI/GLM-5.3-Flash-NVFP4
```

Downloads last month

Safetensors

Model size 169B params

Tensor type F32 · BF16 · U8 · F8_E4M3

 Chat template

 Files info

Inference Providers **NEW**

 Image-Text-to-Text

This model isn't deployed by any Inference Provider.



Ask for provider support

Model tree for RedHatAI/GLM-5.3-Flash

Base model [zai-org/GLM-5.3-Flash](#)
 Quantized (37) [this model](#)

Quantizations    1 model

```
--tensor-parallel-size 4 \  
--no-enable-flashinfer-autotune \  
--tool-call-parser glm47 \  
--enable-auto-tool-choice \  
--reasoning-parser glm45
```

Enable Speculative Decoding

```
docker run --gpus all \  
  --privileged --ipc=host -p 8000:8000 \  
  -v ~/.cache/huggingface:/root/.cache/huggingface \  
  -e VLLM_ENGINE_READY_TIMEOUT_S=3600 \  
  -e PYTORCH_CUDA_ALLOC_CONF=expandable_segments=true \  
  vllm/vllm-openai:glm53-flash RedHatAI/GLM-53-Flash-72B-  
  --tensor-parallel-size 4 \  
  --no-enable-flashinfer-autotune \  
  --tool-call-parser glm47 \  
  --enable-auto-tool-choice \  
  --reasoning-parser glm45 \  
  --gpu-memory-utilization 0.85 \  
  --disable-custom-all-reduce \  
  --speculative-config '{"method":"mtp","num
```

Evaluations

Benchmark	Metric	Avg Score
GPQA Diamond	gpqa_pass@k:k=1 (3 seeds)	90.57%
AIME25	pass@k:k=1&n=1 (8 seeds)	86.67%
GSM8K Platinum CoT	exact_match,strict-match (3 seeds)	97.74%
MATH-500	pass@k:k=1&n=1 (3 seeds)	94.87%



System theme

TOS

Privacy

About

Careers



Models

Datasets

Spaces

Pricing

Docs