

RedHatAI / **Kimi-K2.6-NVFP4** 📄

🤍 like

10

Follow

Red Hat AI

2.44k

🔗 Safetensors

kimi_k25

kimi

nvfp4

vllm

compressed-tensors

custom_code

⋮

📄 Model card

📁 Files

🔗 xet



Community

NVFP4-Quantized RedHatAI/Kimi-K2.6-NVFP4

This is a preliminary version (and subject to change) of NVFP4-quantized [moonshotai/Kimi-K2.6](#) model, for performant inference on NVIDIA Blackwell GPUs. The model has both weights and activations quantized in NVFP4 format with [vllm-project/llm-compressor](#).

It is compatible and tested against vllm v0.20.0. Deploy it via vllm serve using the recipes at <https://recipes.vllm.ai/moonshotai/Kimi-K2.6>.

Creation Script:

Kimi K2.6 support will land in <https://github.com/vllm-project/llm-compressor/pull/2662>. The script to create will be posted as a link shortly

Preliminary Evaluations

1. GSM8K Platinum:

```
lm_eval --model local-chat-completions \
  --tasks gsm8k_platinum_cot_llama \
```

Downloads last
month
8,545



⚡ Inference Providers NEW

This model isn't deployed by any Inference Provider.



1

[Ask for provider support](#)

🔗 Model tree for RedHatAI/Kimi-K2.6

Base model [moonshotai/Kimi-K2.6](#)

🔗 Quantized (30) [this model](#)

```
--model_args "model=RedHatAI/Kimi-K2.6-NVFI
--num_fewshot 0 \
--apply_chat_template \
--gen_kwargs "do_sample=True,temperature=1
```

Recovery:

	moonshot ai/Kimi-K2.6 (original in W4A16)	Red Hat AI/Kimi-K2.6-NVFP4 (t his model)
Accuracy	94.29	93.96
Recovery	-	99.6%

Note: More rigorous evaluations are currently in progress and will be available soon.