

z-lab / **Qwen3.8-27B-DFlash2**   like 224 Follow  Z Lab 1.19k

 Text Generation
  Transformers
  Safetensors
  qwen3
  dflash2
  speculative-decoding

 block-diffusion
  draft-model
  sglang
  vllm
  text-generation-inference
  License: apache-2.0


 Deploy
  Copy to bucket **NEW**
 Use this model

 **Model card**
 Files
  xet
  Community 11

Qwen3.8-27B-DFlash2

[Blog](#) | [GitHub](#)

This repository contains the DFlash 2 draft model for [Qwen/Qwen3.8-27B](#). It is not a standalone language model: it runs inside a speculative decoding server and drafts tokens for the target model to verify. This repository is a mirror of [incoai/Qwen3.8-27B-DFlash2](#).

DFlash 2 is a block-diffusion drafter for speculative decoding. It predicts a whole block of tokens in a single pass and keeps the top candidates at every position. A lightweight selector then traces one coherent path through them. Two-tap dynamic convolutions in the backbone keep the draft from decaying toward the end of the block. Decoding is lossless: greedy output matches the target model exactly, and sampling preserves its distribution.

Downloads last month
64,984



Safetensors

Model size 2B params
 Tensor type BF16  Files info

Inference Providers **NEW**

 Text Generation

This model isn't deployed by any Inference Provider.


 1 Ask for provider support

Model tree for z-lab/Qwen3.8-27B...

Base model [Qwen/Qwen3.8-27B](#)
 Finetuned (197) [this model](#)
 Finetunes 1 model
 Quantizations    8 models

Collection including z-lab/Qwen3...

DFlash 2  Collection



Quick Start

Serve with [SGLang](#):

```
pip install "sglang[all] @ git+https://github.com/sgl-lab/sglang"
```

```
python -m sglang.launch_server \
  --model-path Qwen/Qwen3.8-27B \
  --speculative-algorithm DFLASH \
  --speculative-draft-model-path incoai/Qwen3.8-27B-DFlash2 \
  --speculative-num-draft-tokens 8
```

Or with [vLLM](#):

```
pip install -U "vllm @ git+https://github.com/vllm-project/vllm"
```

```
vllm serve Qwen/Qwen3.8-27B \
  --speculative-config '{
    "method": "dflash",
    "model": "incoai/Qwen3.8-27B-DFlash2",
    "num_speculative_tokens": 7
  }'
```

See the [blog post](#) for other engines and more details.

Evaluation

- Runtime: SGLang on one NVIDIA H200, with FlashAttention 3 for target and draft attention

- Speculation block size: 8 (7 draft tokens per verification step)
- Sampling: Qwen3.8's officially recommended parameters (temperature 1.0, top-p 0.95, top-k 20), with xhigh reasoning effort
- Maximum new tokens: 4096
- Prompts: benchmark formatting from [z-lab/dflash](#)

We compare autoregressive decoding, Qwen3.8's built-in seven-token MTP, a community DSpark drafter ([RadixArk/Qwen3.8-27B-DSpark](#)), and DFlash 2. All speculative methods propose seven draft tokens per verification step.

Acceptance Length

Acceptance length is the per-request mean of completion tokens divided by verification steps. Higher is better.

Task	MTP	DSpark	DFlash 2
GSM8K	5.02	4.36	5.46
MATH-500	4.72	3.92	5.28
HumanEval	3.91	3.30	4.39
MBPP	3.99	3.51	4.79
MT-Bench	3.74	3.01	4.10

Throughput

Throughput is total output tokens divided by end-to-end wall time. Each cell shows output tok/s (speedup vs. autoregressive).

Concurrency 1

Task	Autoregressive	MTP	DSpark	DFlash 2
GSM8K	68.9	178.5 (2.59×)	185.3 (2.69×)	236.1 (3.43×)
MATH-500	69.0	172.8 (2.51×)	174.5 (2.53×)	230.7 (3.34×)
HumanEval	69.0	151.9 (2.20×)	159.9 (2.32×)	214.6 (3.11×)
MBPP	69.0	153.1 (2.22×)	163.3 (2.37×)	226.9 (3.29×)
MT-Bench	68.9	134.9 (1.96×)	137.6 (2.00×)	184.0 (2.67×)

Concurrency 8

Task	Autoregressive	MTP	DSpark	DFlash 2
GSM8K	467.2	1,022.1 (2.19×)	1,040.8 (2.23×)	1,328. (2.84×)
MATH-500	480.0	1,023.5 (2.13×)	1,025.8 (2.14×)	1,368. (2.85×)
HumanEval	483.4	934.2 (1.93×)	956.5 (1.98×)	1,291. (2.67×)

Task	Autoregressive	MTP	DSpark	DFlash 2
MBPP	478.0	938.1 (1.96×)	974.1 (2.04×)	1,328.0 (2.78×)
MT-Bench	480.5	835.2 (1.74×)	802.3 (1.67×)	1,090.0 (2.27×)

Concurrency 32

Task	Autoregressive	MTP	DSpark	DFlash 2
GSM8K	1,329.8	1,381.1 (1.04×)	1,506.5 (1.13×)	1,922.0 (1.45×)
MATH-500	1,505.8	1,415.6 (0.94×)	1,429.0 (0.95×)	1,951.0 (1.30×)
HumanEval	1,546.5	1,296.8 (0.84×)	1,330.1 (0.86×)	1,799.0 (1.16×)
MBPP	1,507.7	1,314.9 (0.87×)	1,361.3 (0.90×)	1,886.0 (1.25×)
MT-Bench	1,507.4	1,159.7 (0.77×)	1,115.5 (0.74×)	1,525.0 (1.01×)

Citation

If you find DFlash 2 useful, please cite:

```
@misc{inco2026dflash2,
  title = {{DFlash 2: Keep Drafting Parallel}},
  author = {{Inco AI}},
  year = {2026},
  month = {August},
```

```
url      = {https://inco.ai/blog/dflash2/}  
}
```

Please also cite the original DFlash paper:

 System theme

[TOS](#)

[Privacy](#)

[About](#)

[Careers](#)



[Models](#)

[Datasets](#)

[Spaces](#)

[Pricing](#)

[Docs](#)

```
@inproceedings{chen2026dflash,  
  title      = {{DFlash: Block Diffusion for I  
  author     = {Chen, Jian and Liang, Yesheng  
  booktitle  = {International Conference on M  
  year       = {2026}  
}
```