



zai-org/GLM-5.2



like

2.17k

Follow



Z.ai

15.3k



Text Generation



Transformers



Safetensors



English



Chinese

glm_moe_dsa

conversational



Eval Results

arxiv:2602.15763

arxiv:2603.12201



License: mit



Deploy

→ Copy to bucket

NEW

Use this model



Model card



Files



xet



Community 20



GLM-5.2

Join our [WeChat](#) or [Discord](#) community.Check out the GLM-5.2 [blog](#) and GLM-5 [Technical report](#).Use GLM-5.2 API services on [Z.ai API Platform](#).Try GLM-5.2 [here](#).[\[Paper\]](#) [\[GitHub\]](#)

Introduction

We're introducing GLM-5.2, our latest flagship model for long-horizon tasks. It marks a substantial leap in long-horizon task capability over its predecessor GLM-5.1 and, for the first time, delivers that capability on a **solid 1M-token context**. GLM-5.2's new capabilities include:

- **Solid 1M Context:** A solid 1M-token context that stably sustains long-horizon work
- **Advanced Coding with Flexible Effort:** Stronger coding capabilities with multiple thinking effort levels to balance performance and latency
- **Improved Architecture:** We propose [IndexShare](#), which reuses the same indexer across every four sparse attention layers, reducing per-token FLOPs by 2.9× at a 1M context length. We also improve GLM-5.2's MTP layer for speculative decoding, increasing the acceptance length by up to 20%

Downloads last

month

40,127



Safetensors

Model size 753B params

Tensor type BF16 · F32

Chat template

Files info

Inference Providers

Text Generation

Input a message
chatting with zai-org/GLM-5.2

Run 15,000+ Models Instantly

Inference Providers let you run inference on thousands of models served by our partners using a simple, unified, OpenAI-compatible serverless API ([Learn more](#)).

zai-org/GLM-5.2 is supported by the following Inference Providers:



Zai



Together AI



Novita



Fireworks



Featherless AI



DeepInfra

View API Code

Dismiss

Your prompt here...

Send

</> View Code

Compare providers

Model tree for zai-org/GLM-5.2

Adapters 4 models

Finetunes 8 models

Quantizations 58 models

Spaces using zai-org/GLM-5.2 35

smolagents/ml-intern

akhaliq/GLM-5.2


angelorovatti/zai-org-GLM-5.2

luizfelipeima/zai-org-GLM-5.2

zoom-ai/hle-leaderboard

- **Pure Open:** An MIT open-source license — no regional limits, technical access without borders

LLM Performance Evaluation



Benchmark

Benchmark	GLM-5.2	GLM-5.1	Qwen3.7-Max	MiniMax M3
Reasoning				
HLE	40.5	31	41.4	37
HLE (w/ Tools)	54.7	52.3	53.5	-
CritPt	20.9	4.6	13.4	3.7
AIME 2026	99.2	95.3	97	-
HMMT Nov. 2025	94.4	94	95	84.4
HMMT Feb. 2026	92.5	82.6	97.1	84.4
IMOAnswerBench	91.0	83.8	90	-
GPQA-Diamond	91.2	86.2	90	93
Coding				
SWE-bench Pro	62.1	58.4	60.6	59
NL2Repo	48.9	42.7	47.2	42.1
DeepSWE	46.2	18	18	20
ProgramBench	63.7	50.9	-	-
Terminal Bench 2.1 (Terminus-2)	81.0	63.5	75	65
Terminal Bench 2.1 (Best Reported Harness)	82.7	69	-	-
FrontierSWE (Dominance)	74.4	30.5	-	-

grafvbn1212/dijfsuiofo

grafvbn1212/API-Krodum

SherlockDev9813/API-Krodum

+ 27 Spaces

Collection including zai-org/GLM-...

GLM-5.2 Collection

2 items • Updated 7 days ago • 44

Papers for zai-org/GLM-5.2

IndexCache: Accelerating Sparse At...

Paper • 2603.12201 • Publi... • 60

GLM-5: from Vibe Coding to Agentic...

Paper • 2602.15763 • Pub... • 180

Article mentioning zai-org/GLM-5.2

GLM-5.2: Built for Long-Horizon Tasks

zai-org • 6 days ago • 99

Evaluation results

ldavidrein/gpqa · Diamond 91.2

ScaleAI/SWE-bench_Pro ... 62.1

crosbylegal/RedlineBen... 45.7*

datacurve/deep-swe · D... 46.2

cais/hle · Hle

default 40.5

With tools 54.7*

Benchmark	GLM-5.2	GLM-5.1	Qwen3.7-Max	MiniMax M3
PostTrainBench	34.3	20.1	-	-
SWE-Marathon	13.0	1.0	-	-
Agentic				
MCP-Atlas (Public Set)	76.8	71.8	76.4	74.2
Tool-Decathlon	48.2	40.7	-	-

Serve GLM-5.2 Locally

GLM-5.2 supports deployment with the following frameworks. Feel free to try them out:

- [SGLang](#) (v0.5.13.post1+) — see [cookbook](#)
- [vLLM](#) (v0.23.0+) — see [recipes](#)
- [Transformers](#) (v0.5.12+) — see [transformers docs](#)
- [KTransformers](#) (v0.5.12+) — see [tutorial](#)
- [Unsloth](#) (v0.1.47-beta+) — see [guide](#)
- For deployment on the Ascend NPU platform, inference frameworks such as vLLM-Ascend, xLLM and SGLang are supported — see [here](#).

Footnote

- **Humanity’s Last Exam (HLE) & other reasoning tasks:** We use sampling parameters of temperature=1.0, top_p=0.95 for evaluation. We evaluate with a maximum generation length of 163,840 tokens. By default, we report the text-only subset; results marked with * are from the full set. For AIME, HMMT and IMOAnswerBench, we evaluate each question using the following system prompt:
Your response should be in the following format:\nExplanation: {your explanation for your final answer}\nExact Answer: {your succinct, final answer}\nConfidence: {your confidence score between 0% and 100% for your answer}. We use GPT-5.5 (medium) as the judge model. For HLE-with-tools, we use

a maximum context length of 300,000 tokens, with no context management strategy.

- **SWE-Bench Pro:** We run the SWE-Bench Pro suite with OpenHands using a tailored instruction prompt. Settings: `temperature=1`, `top_p=1`, `max_new_tokens=32k`, with a 400K context window.
- **NL2Repo:** We evaluated NL2Repo with `temperature=1.0`, `top_p=1.0`, and `max_new_tokens=48k` under 400k context. To prevent hacking, we use rule-based and a LLM-based judgement to prevent malicious behaviors (e.g., unauthorized pip or curl operations).
- **DeepSWE:** We run DeepSWE with the official pier evaluation framework and the mini-swe-agent harness (`temperature=1.0`, `top_p=1.0`, `timeout=2h`, 400K context). Each task is solved in an isolated container with 2 CPUs, 8 GB RAM, and no internet access.
- **ProgramBench:** We evaluate ProgramBench (200 instances) with Claude-Code 2.1.156 using `temperature=1.0`, `top_p=1.0`, `max_tokens=64000`, `max_turns=2000`, `sample_timeout=6h`, `reasoning_effort=max`, with a 400K context window. Each instance runs in a (4 CPUs, 8 GB RAM) sandbox with internet access disabled.
- **Terminal-Bench 2.1 (Terminus 2):** We evaluate Terminal-Bench 2.1 with Terminus-2 framework using `parser=json`, `timeout=4h`, `temperature=1.0`, `top_p=1.0`, `max_new_tokens=48k`, `max_episodes=500`, with a 256K context window. Resource limits are capped at 4 CPUs and 8 GB RAM.
- **Terminal-Bench 2.1 (Claude Code):** We evaluate in Claude Code 2.1.167 with `temperature=1.0`, `top_p=0.95`, `max_new_tokens=131072`. We override `max_new_tokens` to 128k via a transparent proxy, bypassing the 64k CLI cap to restore the configurability of `CLAUDE_CODE_MAX_OUTPUT_TOKENS`. We remove wall-clock time limits, while preserving per-task CPU and memory constraints. Scores are averaged over 5 runs.

- **MCP-Atlas:** All models were evaluated in think mode on the 500-task public subset with a 10-minute timeout per task. We use Gemini-3.0-Pro as the judge model for evaluation.
- **Tool-Decathlon:** We use the official evaluation service and set max_token to 128K.
- **FrontierSWE:** The evaluation was conducted by [Proximal](#) with 1M context length, max effort level, and 128K maximum output tokens. Dominance score reported as of 2026/06/16.
- **PostTrainBench:** The evaluation was conducted by [PostTrainBench](#) with 1M context length, max effort level, and 128K maximum output tokens.
- **SWE-Marathon:** The evaluation was conducted by [Abundant AI](#) with 1M context length, max effort level, and 128K maximum output tokens.

Citation

If you find GLM-5.2 useful in your research, please cite our technical report:

```
@misc{glm5team2026glm5vibecodingagentic,
  title={GLM-5: from Vibe Coding to Agent},
  author={GLM-5-Team and : and Aohan Zeng},
  year={2026},
  eprint={2602.15763},
  archivePrefix={arXiv},
  primaryClass={cs.LG},
  url={https://arxiv.org/abs/2602.15763}
}
```